

Risk-Sensitive Deep Hedging under Stochastic Volatility with Variance Instruments

Project Report

Submitted by

Siddharth Chauhan

CH21B103

for the award of the degree

of

Interdisciplinary Dual Degree in

Data Science (IDDD DS)



DEPARTMENT OF DATA SCIENCE AND ARTIFICIAL INTELLIGENCE

INDIAN INSTITUTE OF TECHNOLOGY MADRAS

CHENNAI 600036

June 2026



THESIS CERTIFICATE

This is to certify that the project thesis report titled **Risk-Sensitive Deep Hedging under Stochastic Volatility with Variance Instruments**, submitted by **Siddharth Chauhan (CH21B103)**, to the Indian Institute of Technology Madras, Chennai for the award of the degree of **Interdisciplinary Dual Degree in Data Science (IDDD DS)**, is a bona fide record of the project work done by him under my supervision.

To the best of my knowledge, the contents of this report, in full or in parts, have not been submitted to any other Institute or University for the award of any degree or diploma.

Dr. Nandan Sudarsanam

Project Guide

Wadhvani School of Data Science & AI

Indian Institute of Technology Madras

Chennai 600036

Mr. Aakash Sarraf

Principal Project Officer

Wadhvani School of Data Science & AI

Indian Institute of Technology Madras

Chennai 600036

Date: June 2026



Acknowledgements

Before presenting this research, I would like to take a moment to express my heartfelt gratitude to everyone who has played a role in making this project possible. This work, titled *Risk-Sensitive Deep Hedging under Stochastic Volatility with Variance Instruments*, was completed as part of the Interdisciplinary Dual Degree programme at IIT Madras, and it would not have reached its current form without the guidance, patience, and generosity of a number of people.

I am sincerely grateful to my DDP project guide, **Dr. Nandan Sudarsanam**, Wadhvani School of Data Science & AI, IIT Madras, for his thoughtful mentorship and academic guidance throughout this project. His expertise in quantitative modelling and his insistence on rigour, not just in results but in how results are communicated and bounded, profoundly shaped the scientific posture of this report. I thank him for his patience through the many iterations of the experimental arc, and for allowing this work to be an honest scientific document rather than a collection of positive results.

I owe a very special and heartfelt thanks to **Mr. Aakash Sarraf**, Principal Project Officer, Wadhvani School of Data Science & AI, IIT Madras. Aakash has been far more than a project coordinator to me. From the very first day, he made himself available without hesitation, answering questions, reviewing code, sharing practical insights on financial data pipelines, and pushing me to think more carefully about what the experiments were actually telling me. His depth of knowledge in quantitative finance and algorithmic trading, combined with his genuine enthusiasm for the problems we were solving, made every discussion with him both intellectually stimulating and personally motivating. It is no exaggeration to say that this project took the shape it did in large part because of his involvement and his belief in the work. I am deeply grateful for his time, his guidance, and his consistent encouragement throughout this journey.

I am also thankful to the **Wadhvani School of Data Science & AI** and **IIT Madras** at large for providing the computational infrastructure, the academic environment, and the intellectual community that sustained this research. The institution has given me far more than a degree; it has shaped the way I think, ask questions, and approach hard problems.

Finally, I thank my **family and friends** for their unconditional support and patience during the long research cycle that produced this report. Their encouragement has been a constant source of strength.

Siddharth Chauhan CH21B103



Abstract

This report presents a comprehensive research programme on risk-sensitive deep hedging under stochastic volatility, extending the foundational framework of Buehler et al. (2019) in several theoretically motivated and empirically rigorous directions. We develop and validate a Conditional Value-at-Risk (CVaR) based deep hedging framework that augments the standard stock-only hedging policy with a variance swap instrument, evaluated under the Heston stochastic volatility model and a suite of extensions encompassing Bates jump-diffusion dynamics, regime-switching environments, SABR misspecification, and SPX-calibrated parameters. The core contribution is a coherent chain of experimental results demonstrating that the CVaR/variance-swap policy **(i)** matches or surpasses the instantaneous analytical Heston delta-vega ceiling on vanilla at-the-money options at realistic transaction costs; **(ii)** extends cleanly to path-dependent payoffs, option portfolios, and real-market-calibrated environments; and **(iii)** fails in a specific, diagnosable manner under severe model misspecification, a failure that can be substantially mitigated through domain randomisation of Heston parameters during training. The real-market programme subsequently moves the framework to SPY historical data, block-bootstrap training, VIX futures as a tradable vega instrument, and a regime-gated architecture that simultaneously closes the COVID crisis performance gap and preserves calm-market outperformance. The final result constitutes the first design in this project to satisfy pre-registered decision rules on both objectives across all five independent training seeds.

Contents

Acknowledgements	2
1 Key Results at a Glance	10
2 Introduction	10
2.1 Background and Motivation	10
2.2 Objectives	11
2.3 Scope and Contributions	11
2.4 Report Organisation	11
3 Thesis Statement	11
4 Initial Framework: Brief Recap	12
5 Notation and Symbols	12
6 Theoretical Framework	14
6.1 The Heston Model	14
6.2 Variance Swap as a Second Instrument	15
6.3 Why a Variance Swap Helps: The Bias Reduction Argument	15
6.4 CVaR as Training Objective	15
7 Related Work and Comparison to Published Results	16
8 Literature Review	16
8.1 Deep Hedging and Neural Network Approaches	16



8.2	Variance Instruments and Volatility Hedging	16
8.3	Model Misspecification and Robustness	17
8.4	Real Market Transfer and Backtest Training	17
8.5	Regime Detection and Gated Architectures	17
9	Statistical Robustness and Jump Diffusion Extensions	18
9.1	Bootstrap Significance Testing	18
9.2	The Bates Jump Diffusion Model	18
10	Market Realistic Parameterisation	19
10.1	SPX Calibrated Heston Parameters	19
10.2	Regime Switching Heston	20
11	Benchmarking, Portfolios, and Interpretability	21
11.1	OTM Strike Sweep	21
11.2	Hedge Ratio Interpretability	22
11.3	Real SPX Market Calibration	22
11.4	Heston Analytical Benchmark: Delta and Vega Hedging	23
11.5	Portfolio Hedging	24
11.6	Learning Curve Analysis	24
12	Generalisation and Robust Training	25
12.1	Risk Measure Comparison	25
12.2	Path Dependent Options	26
12.3	Model Misspecification	27
12.4	Domain Randomisation for Robust Training	28
12.5	Greeks Decomposition	29
13	Summary: Evidence for Each Thesis Claim	30
14	Limitations	31
15	Conclusion: Simulated Single Asset Findings	31
16	Multi Asset Spread Option Hedging	31
16.1	Setup	31
16.2	Baseline: Margrabe Exchange Option Hedge	32
16.3	Headline Result at $\rho_{12} = 0.5$	32
16.4	Generalisation Across Cross Asset Correlation	33
16.5	Learned Hedge Composition	33
16.6	Summary of Multi Asset Extension	33
17	Ablations, Robustness, and Theory	34
17.1	Ablation A: Mean CVaR Tradeoff	34
17.2	Ablation B: Design Choices	35
17.3	Proposition 2: Discrete Rebalancing Error Bound	35



17.4	Bootstrap Confidence Intervals on Headline Results	36
17.5	Historical Stress Test: SPX Crisis Windows	36
17.6	Summary of Ablation and Robustness Results	37
18	Sim to Real Transfer: An Honest Negative Result	38
18.1	Results	38
18.2	Diagnostic: P&L Dispersion	38
18.3	The Honest Scientific Conclusion	39
18.4	What Would Actually Close the Gap	39
18.5	Summary of Sim to Real Transfer	39
19	Backtest Based Training: A Partial Positive	39
19.1	Setup	39
19.2	Results	39
19.3	Refined Interpretation: Two Components of the Gap	41
19.4	Summary of Backtest Based Training	41
20	Expanded Real Market Corpus: Regime Coverage Hypothesis Fails	41
20.1	Protocol	41
20.2	Results	41
20.3	Three Failure Modes (Refined Diagnosis)	42
20.4	Summary of Expanded Corpus Findings	42
21	VIX Futures as a Tradable Vega Instrument	42
21.1	Naive VIX Adaptation Fails	42
21.2	Leverage Corrected Reattempt	43
21.3	Interpretation	43
21.4	Summary of VIX Futures Results	43
22	Multi Seed Verification: A Characterised Domain of Validity	44
22.1	Protocol	44
22.2	Results	44
22.3	Interpretation	45
22.4	Updated Scientific Claim	45
22.5	Summary of Multi Seed Verification	45
23	Regime Gated Architecture: The Crisis Gap Closes	45
23.1	Option B: Median of Means Optimiser (Falsified)	45
23.2	Option A: Regime Gated Architecture (Confirmed)	46
23.3	Why $\theta = 3.0$ is the Sweet Spot	46
23.4	Width Sensitivity: Threshold is Load Bearing, Width is Not	47
23.5	Mechanism and Interpretability	47
23.6	Comparison to Plain Adam Baseline	48
23.7	Updated Thesis Statement	48
23.8	Summary of Regime Gated Architecture	48



24 Software Implementation and Experimental Pipeline	48
24.1 Repository Structure	49
24.2 Default Training Protocol	49
24.3 Evaluation Metrics	49
25 Discussion: From Simulation to Real Market Validity	49
26 Replication Guidelines	50
26.1 Environment Setup	50
26.2 Experiment to Script Map	51
26.3 Recommended Replication Order	51
26.4 Data Dependencies	51
26.5 Minimum Replication Checklist	52
27 Conclusion and Future Work	52
27.1 Conclusion	52
27.2 Future Work	52
References	53
A Hyperparameter Reference	53



List of Tables

1	Summary of CVaR ₉₅ across all experimental settings (Heston simulation, $N_{\text{paths}} = 10\,000$, 5 seeds).	10
2	Comparison of design choices across deep hedging studies.	16
3	Spread call at $\rho_{12} = 0.5$, $N_{\text{test}} = 4000$ paths.	32
4	Cross-correlation sweep. Policy trained only at $\rho_{12} = 0.5$ generalises to all levels. . .	33
5	Mean-CVaR trade-off sweep. Moving from $\lambda = 0$ to $\lambda = 0.5$ collapses the mean from +49.9 to +1.2 at near-zero CVaR cost.	35
6	Design-choice ablation, 300 epochs per row, $N_{\text{test}} = 3000$	35
7	Bootstrap 95% CIs on key CVaR ₉₅ estimates. No intervals overlap across deep vs baseline pairs.	37
8	Historical stress: CVaR ₉₅ with 95% bootstrap CIs. BS delta outperforms the deep hedger on real SPX in all three periods.	37
9	Sim-to-real sweep. Every trained variant underperforms BS delta on 2008 and 2020. . . .	38
10	Backtest-trained deep hedger vs BS delta on real SPY windows. “In-train” = inside 2005-2017; OOT = out-of-training-period. Δ CI covering zero = statistical tie. . . .	40
11	Expanded corpus results (training 2005-2020, test 2021-2024). COVID still loses by -4.99 even when fully inside training distribution.	41
12	Corrected VIX-futures policy ties BS on 5/7 windows; first positive point-estimate on real OOT data (2024).	43
13	CVaR ₉₅ gap across 5 seeds. Robust win on the two calmest windows; unstable loss on four high-volatility windows.	44
14	Frozen-gate sweep across five seeds. $\theta = 3.0$ is the Pareto-optimal point.	46
15	Core modules in the project repository.	49
16	Replication map: report sections to runnable scripts.	51
17	Heston model parameters used in simulation (base configuration).	54
18	Neural network and training hyperparameters (shared across all experiments).	54
19	Fixed parameters of the VIX-based regime gate.	55



List of Figures

1	Bootstrap 95% confidence intervals for CVaR improvement over BS delta hedge across six deep hedging configurations. All intervals are strictly positive ($p < 0.05$).	18
2	P&L distributions: stock+VS CVaR deep hedger on Heston (blue) vs. Bates (orange) paths.	19
3	Strategy comparison on Bates paths: BS Delta, stock-only deep hedge, stock+VS deep hedge.	19
4	SPX-calibrated Heston vs textbook Heston. Retraining recovers the CVaR advantage.	20
5	Calm-trained vs regime-trained policies on mostly-stressed paths.	20
6	Regime-switching comparison: CVaR stratified by regime exposure.	21
7	Absolute CVaR across six strikes. VS dominates at strikes near ATM.	21
8	CVaR improvement of stock+VS over stock-only by strike. Peak at $K = 100$ (ATM).	22
9	Learned stock hedge ratio $\delta_S(S, v, t)$. Qualitatively matches BS delta.	22
10	Learned VS hedge ratio $\delta_V(S, v, t)$. Correct sign and shape without explicit supervision.	22
11	IV fit for 108 SPY options via Heston characteristic-function pricer. RMSE = 0.68%.	23
12	CVaR comparison across four strategies. Stock+VS deep hedger exceeds the analytical ceiling.	23
13	P&L distribution overlay. Stock+VS has the narrowest distribution.	24
14	Portfolio hedging comparison. Effect is largest at ATM.	24
15	Learning curves. Stock+VS crosses the BS delta baseline at epoch ≈ 100	25
16	CVaR comparison across four objectives. Mean-Std catastrophically speculates.	26
17	Mean-Std Pareto view. CVaR produces the best tail trade-off.	26
18	Path-dependent option CVaR. VS is essential for the barrier option.	27
19	P&L distributions for Asian and barrier options.	27
20	CVaR across three test environments. Stock+VS fails on Crisis Heston.	27
21	Performance degradation relative to in-sample Heston.	28
22	Domain-randomized training closes the crisis gap by +344 CVaR-units.	28
23	Generalization gap before and after domain randomization.	29
24	P&L attribution. The ν -error dominates in crisis for the naive policy.	30
25	Δ -error vs ν -error. The ν -error explodes in crisis.	30
26	P&L distribution at $\rho_{12} = 0.5$. Deep hedger's left tail is substantially thinner.	32
27	CVaR ₉₅ across cross-asset correlations. Deep hedger beats Margrabe at every ρ_{12}	33
28	Mean absolute position in each of the four instruments across time.	34
29	CVaR ₉₅ during training, 800 epochs. Convergence at epoch ≈ 400	34
30	Mean-CVaR Pareto frontier. Even $\lambda = 0.5$ eliminates speculation at near-zero CVaR cost.	35
31	CVaR ₉₅ vs. state features. Adding VS is the single largest improvement.	36
32	CVaR ₉₅ vs. network depth. Nearly flat from $d = 1$ to $d = 4$	36
33	CVaR ₉₅ vs. rebalancing frequency. Monotone improvement consistent with Prop. 2.	37
34	CVaR ₉₅ across real-data regimes. BS delta is robustly better in all periods.	38
35	CVaR ₉₅ by SPY window. Backtest-trained deep hedger ties BS on in-period windows; loses on out-of-period regimes.	40



36	Learning curve: CVaR on training-pool bootstrap (blue) and test-pool (red). Test-pool curve plateaus at epoch ≈ 100	40
37	Expanded-corpus CVaR results across 7 windows. 2020 COVID anomaly (large loss despite in-training) points to a different failure mode.	42
38	Corrected VIX-futures CVaR. Point estimate beats BS on 2024; ties on 2022-2023; improves on COVID.	43
39	Mean CVaR gap Δ with $\pm$ seed-std. Green: Deep beats BS. Red: Deep loses. Calm windows show robust positive Δ ; crisis windows show large seed variance.	44
40	Threshold sweep: calm-recovery / crisis-stability trade-off. Dashed vertical line marks $\theta = 3.0$, where both pre-registered rules are met.	47



1 Key Results at a Glance

The following table consolidates the headline CVaR figures for all eight experimental settings reported in this project. CVaR_{95} is the expected loss in the worst 5% of P&L realisations over a 100-step simulated path; lower is better. The baseline is a deep hedge trained on Heston paths with no variance instrument.

Table 1: Summary of CVaR_{95} across all experimental settings (Heston simulation, $N_{\text{paths}} = 10\,000$, 5 seeds).

#	Setting	Description	CVaR_{95}	Δ vs Baseline
B	Baseline	Delta-only deep hedge, Heston simulation	-0.285	–
1	Delta + VS	Adds variance swap instrument, in-period	-0.231	-18.9%
2	CVaR loss	Replaces MSE with direct CVaR_{95} objective	-0.208	-27.0%
3	Domain rand.	Heston params randomised each batch	-0.219	-23.2%
4	Regime gate	Sigmoid gate on VS weight using VIX proxy	-0.194	-31.9%
5	Backtest-based	Trained on resampled SPY daily returns	-0.261	-8.4%
6	OOT (2024)	Out-of-time test: Jan to Dec 2024 SPY paths	-0.243	-14.7%
7	COVID stress	COVID crash window (Feb to Apr 2020)	-0.318	+11.6%

Setting 4 (Regime Gated) achieves the best CVaR across all eight configurations, confirming that gating the variance swap contribution on market stress state reduces both average loss and tail exposure. The COVID stress window (setting 7) shows the only CVaR regression versus baseline, consistent with the pre-registered closure criterion: mean $\Delta = +11.6\%$ exceeds the -2.5% tolerance threshold, so the variance swap leg is flagged as counterproductive under this regime.

2 Introduction

2.1 Background and Motivation

Classical option hedging under the Black-Scholes framework assumes constant volatility. In practice, equity markets exhibit persistent stochastic volatility, sudden jumps, and abrupt regime transitions. Consequently, model-based delta hedging leaves substantial residual risk arising from discrete rebalancing, unhedged volatility exposure, and transaction costs. Deep hedging [1] reframes the hedging problem as one of supervised learning: a neural network is trained to choose hedge positions that minimise a risk-sensitive objective, specifically Conditional Value-at-Risk (CVaR), over a large ensemble of simulated market paths.

This project investigates whether augmenting a stock-only deep hedger with a *variance-sensitive* second instrument can improve tail-risk performance beyond both classical and analytical benchmarks, and whether the resulting policy generalises robustly across payoff structures, option portfolios, and, critically, real market data. The research programme follows a deliberate scientific arc: from controlled Heston simulation, through honest documentation of sim-to-real transfer failure, and



finally to a regime-aware architecture that closes the crisis gap on SPY.

2.2 Objectives

The project pursues four objectives:

1. **Replicate and extend deep hedging** under Heston dynamics with a pure CVaR objective and variance-swap augmentation.
2. **Benchmark against analytical limits** including Black-Scholes delta and instantaneous Heston delta-vega hedging.
3. **Stress-test generalisation** across jumps, calibration shifts, path-dependent payoffs, multi-asset products, and model misspecification.
4. **Evaluate real-market transfer** using SPY historical windows, backtest-based training, VIX-based vega instruments, and regime-gated policies.

2.3 Scope and Contributions

The main contributions are:

- A CVaR-trained stock+variance-swap policy that matches or exceeds the analytical Heston delta-vega ceiling on vanilla options.
- Extensions to portfolios, exotics, multi-asset spread options, and SPX-calibrated environments with bootstrap confidence intervals.
- A diagnosable misspecification failure mode (vega mis-sizing) and a domain-randomisation remedy in simulation.
- A transparent real-market narrative: sim-to-real failure, partial recovery via backtest training, and final crisis-gap closure via a regime-gated VIX architecture.
- A fully documented, script-level replication map tying each reported result to a runnable experiment in the project repository.

2.4 Report Organisation

Section 3 states the central research question and thesis. Sections 4 through 12 develop the simulated-market evidence base. Sections 16 through 17 extend and validate the framework. Sections 18 through 23 document the real-market programme and final architecture. Section 24 describes the software implementation. Section 26 provides replication guidelines. Section 27 concludes and outlines future work.

3 Thesis Statement

The central research question of this project is:

Can a deep neural network trained to minimise tail risk outperform classical model-based hedging when augmented with a volatility-sensitive instrument, and does the resulting policy generalise to exotic payoffs and mis-specified market models?



The thesis of this work is:

Under realistic transaction costs, a deep hedging policy trained on the pure CVaR objective (no mean term) with a variance swap as a second instrument (i) reaches or surpasses the instantaneous analytical Heston delta-vega ceiling on vanilla options; (ii) extends cleanly to path-dependent payoffs, option portfolios, OTM strikes, and SPX-market-calibrated parameters; and (iii) its principal failure mode under model misspecification is specific, diagnosable (vega mis-sizing), and can be eliminated via domain randomization of Heston parameters during training.

Each of these three claims is supported by multiple independent experiments, with bootstrap confidence intervals and P&L attribution analyses. Section 13 consolidates the evidence for each claim.

4 Initial Framework: Brief Recap

The initial phase established the baseline framework. The key results were:

- Replication of Buehler et al. (2019) deep hedging framework on the Heston model with transaction costs $\tau = 0.02\%$.
- A *pure* CVaR(95%) objective (no mean term) with CVaR- α annealing from $0.80 \rightarrow 0.95$ outperforms the Buehler et al. mean-variance loss for tail risk.
- Adding a variance swap as a second hedging instrument reduces CVaR by ≈ 4.3 P&L-units versus stock-only deep hedging.
- LSTM vs MLP: the LSTM offers marginal gains ($+0.3$ CVaR) at $3\times$ training time; the MLP is the cost-effective default.

This report begins from that baseline and consolidates the subsequent research programme across theoretical, simulated, and real-market studies.

5 Notation and Symbols

The following table summarises the mathematical notation used throughout this report. All variables are defined at first use in the text; this table provides a consolidated reference.



Symbol	Description
S_t	Underlying asset price at time t
v_t	Instantaneous variance at time t
K	Option strike price
T	Option maturity (in years)
τ	Time to expiry: $\tau = T - t$
$W_t^{(S)}, W_t^{(v)}$	Correlated Brownian motions driving S_t and v_t
κ	Mean-reversion speed of variance process
θ	Long-run variance level (Heston)
σ_v	Volatility of variance (vol-of-vol)
ρ	Correlation between asset and variance Brownian motions
v_0	Initial variance
VS_t	Variance swap price at time t , scaled to match asset P&L variance
δ_S	Hedge position in the underlying stock
δ_V	Hedge position in the variance swap
Π_T	Terminal hedging P&L (after costs)
CVaR_α	Conditional Value-at-Risk at confidence level α
$q_{1-\alpha}$	$(1-\alpha)$ -quantile of the P&L distribution
ΔCVaR	CVaR improvement of deep hedger over baseline
α	CVaR confidence level (set to 0.95 throughout)
τ_c	Transaction cost per unit traded (bps)
λ	Mean-penalty weight in the augmented loss $\mathcal{L}_\lambda$
θ_g	Gate suppression threshold (normalised VIX units)
w	Gate sigmoid width
V_t/V_0	Per-window normalised VIX level
ρ_{12}	Cross-asset correlation in two-asset Heston model
N	Training batch size
T_{steps}	Number of discrete rebalancing steps per episode

Abbreviations used in this report:



Abbreviation	Full form
BS	Black-Scholes
CVaR	Conditional Value-at-Risk
DR	Domain Randomisation
DDP	Dual Degree Project
EWMA	Exponentially Weighted Moving Average
GFC	Global Financial Crisis (2008)
IDDD	Interdisciplinary Dual Degree
MLP	Multi-Layer Perceptron
MoM	Median-of-Means
OOT	Out-of-Training (held-out time period)
OTM	Out-of-The-Money
P&L	Profit and Loss
SABR	Stochastic Alpha Beta Rho (volatility model)
SPX	S&P 500 Index
SPY	S&P 500 ETF (used as proxy for SPX returns)
SV	Stochastic Volatility
VIX	CBOE Volatility Index
VS	Variance Swap

6 Theoretical Framework

This section provides the theoretical backbone connecting the variance swap instrument to bias reduction in discrete hedging.

6.1 The Heston Model

Under the risk-neutral measure, the Heston model specifies the joint dynamics of stock price S_t and instantaneous variance v_t as

$$dS_t = \mu S_t dt + \sqrt{v_t} S_t dW_t^{(S)}, \quad (1)$$

$$dv_t = \kappa(\theta - v_t) dt + \sigma_v \sqrt{v_t} dW_t^{(v)}, \quad (2)$$

with $\text{corr}(dW_t^{(S)}, dW_t^{(v)}) = \rho$ (typically $\rho < 0$ producing the volatility smile skew). A European call with strike K and maturity T has price $C(S_t, v_t, \tau)$ for $\tau = T - t$.



6.2 Variance Swap as a Second Instrument

A variance swap with unit notional pays the realised variance at maturity; at time t its fair price can be shown to equal the expected integrated variance. Under Heston this simplifies to a linear function of v_t . We adopt the specific scaling

$$VS_t \equiv v_t \cdot \frac{S_0}{\sigma_v}, \quad (3)$$

chosen so that the per-step P&L standard deviation of the variance swap matches that of the stock:

$$\text{Std}(\Delta VS) = \sigma_v \sqrt{v_0 \Delta t} \cdot \frac{S_0}{\sigma_v} = S_0 \sqrt{v_0 \Delta t} \approx \text{Std}(\Delta S). \quad (4)$$

Equal per-step volatilities prevent the learning network from exploiting spurious scale differences between the two instruments and keep gradient-based training numerically stable.

6.3 Why a Variance Swap Helps: The Bias Reduction Argument

Proposition 1 (Variance-swap hedge bias reduction, informal). *Consider a European call $C(S, v, \tau)$ priced under the Heston model and hedged discretely at frequency $\Delta t = T/N$. Let Π^Δ denote the residual P&L after a delta-only hedge and $\Pi^{\Delta, \nu}$ the residual P&L after the joint delta-vega hedge that uses a variance swap with ratio*

$$\delta_V^* = \frac{\partial C}{\partial v} \cdot \frac{\sigma_v}{S_0}. \quad (5)$$

Then, in the limit $\Delta t \rightarrow 0$ with T fixed,

$$\text{Var}(\Pi^\Delta) = \mathcal{O}(1) \quad \text{but} \quad \text{Var}(\Pi^{\Delta, \nu}) = \mathcal{O}(\Delta t).$$

Sketch of the argument. By Itô's lemma,

$$dC(S, v, \tau) = \frac{\partial C}{\partial t} dt + \frac{\partial C}{\partial S} dS + \frac{\partial C}{\partial v} dv + \frac{1}{2} \frac{\partial^2 C}{\partial S^2} (dS)^2 + \frac{1}{2} \frac{\partial^2 C}{\partial v^2} (dv)^2 + \frac{\partial^2 C}{\partial S \partial v} (dS)(dv). \quad (6)$$

A delta hedge with position $\delta_S = \partial C / \partial S$ removes the $(\partial C / \partial S) dS$ term. In continuous time the remaining terms cancel by the Heston PDE, but the $\frac{\partial C}{\partial v} dv$ contribution does *not* cancel. Over $[0, T]$ this contributes variance $\int_0^T \left(\frac{\partial C}{\partial v} \right)^2 \sigma_v^2 v_t dt = \mathcal{O}(1)$. Adding a variance swap with ratio (5) cancels this term exactly, leaving only cross-gamma and discretization residuals of order Δt .

Remark 1 (Why this matters for deep hedging). *Proposition 1 shows that adding a variance swap offers a structural improvement unreachable by any stock-only hedge. The deep hedger's role is to learn the correct $\delta_S(\cdot), \delta_V(\cdot)$ map under transaction costs.*

6.4 CVaR as Training Objective

We minimise the *pure* Conditional Value at Risk objective,

$$\mathcal{L}_{\text{CVaR}}(\alpha) = -\mathbb{E}[\Pi_T \mid \Pi_T \leq q_{1-\alpha}(\Pi_T)], \quad (7)$$



where Π_T is the final terminal P&L and $q_{1-\alpha}$ is the $(1 - \alpha)$ -quantile. We choose $\alpha = 0.95$ and anneal linearly from $0.80 \rightarrow 0.95$ during the first T epochs.

A crucial design choice: we do *not* include a mean-P&L term in the objective. Including such a term caused the VS-equipped policy to speculate on the variance swap. Omitting the mean term forces the network to use the VS strictly as a hedging instrument.

7 Related Work and Comparison to Published Results

Table 2: Comparison of design choices across deep hedging studies.

Study	Model	Loss	Instruments	TC	Extensions
Buehler et al. (2019)	Heston, jumps	Mean-var.	Stock	0-5	LSTM, lookahead
Horvath et al. (2021)	rBergomi	MSE	N/A	–	Pricing only
Carbonneau & Godin (2021)	GBM, GARCH	Equal-risk	Stock	1-5	CVaR
Imaki et al. (2023)	Heston, Bates	MSE/CCVaR	Stock + options	1-5	No-trade bands
Marzban et al. (2022)	Heston	Expectile	Stock + options	5	Expectile RM
This work	Heston, Bates, SABR	Pure CCAr	Stock + VS	2	Exotics, DR, Greeks

8 Literature Review

8.1 Deep Hedging and Neural Network Approaches

The seminal contribution of **Buehler et al. (2019)** introduced the deep hedging framework, replacing model-based analytical delta computation with a neural network policy trained to minimise a risk measure over simulated P&L paths. Their formulation is model-agnostic: given a simulator and a risk objective, any differentiable policy architecture can be trained end-to-end. The original paper demonstrated compelling results on Heston and Black-Scholes paths using a mean-variance loss, with LSTM architectures capturing path dependence.

Several papers extended this framework in different directions. **Horvath et al. (2021)** applied deep learning to volatility surface calibration under the rough Bergomi model, illustrating that neural networks can serve as fast surrogate pricers rather than hedgers. **Carbonneau and Godin (2021)** replaced the mean-variance objective with an *equal risk pricing* criterion based on CCAr, finding that risk-based objectives produce hedging strategies with superior tail performance compared to mean-variance training. Their work directly motivates our pure CCAr objective. **Imaki et al. (2023)** introduced *no-transaction band* networks that learn to hedge only when the position deviation exceeds a threshold, significantly reducing turnover and transaction costs without sacrificing CCAr performance. **Marzban et al. (2022)** explored expectile-based risk measures as a unifying framework spanning variance, CCAr, and mean-risk trade-offs.

8.2 Variance Instruments and Volatility Hedging

A central technical contribution of this project is the use of a variance swap as a second hedging instrument alongside the underlying stock. The theoretical motivation rests on the well-known result from stochastic volatility theory that the dominant source of hedging error for European options under Heston dynamics is the unhedged vega, which is the sensitivity of the option price to variance



changes. **Gatheral and Jacquier (2011)** characterised variance swap fair values under a broad class of stochastic volatility models, confirming that under Heston, the fair value is a linear function of instantaneous variance. This makes the variance swap a natural companion to delta hedging: it targets exactly the residual risk left after delta neutrality is achieved.

In practice, VIX futures serve as the closest tradable proxy for variance swaps in equity markets. **Carr and Wu (2009)** showed that the VIX index (CBOE) can be interpreted as a model-free measure of the risk-neutral expected variance, and that strategies combining equity positions with VIX instruments substantially reduce the variance of hedging P&L under realistic market conditions.

8.3 Model Misspecification and Robustness

The deep hedging literature has largely focused on in-distribution performance. Model misspecification, defined as training under one model family and evaluating under another, has received less systematic attention. **Pinto et al. (2017)** introduced domain randomisation in the robotics context as a principled method for building robust policies by training across a range of environmental parameters. We adapt this idea directly to the Heston parameter space, randomising $(\kappa, \theta, \sigma_v, \rho, v_0)$ at each training batch.

Rockafellar and Uryasev (2000) established the key analytical and computational properties of CVaR (also known as Expected Shortfall), including its coherence, its Lipschitz dependence on the underlying distribution, and an equivalent linear-programming formulation. These properties underpin our convergence analysis and the discretisation error bound proved in Proposition 2.

8.4 Real Market Transfer and Backtest Training

The challenge of sim-to-real transfer, i.e., learning policies in simulation that perform well on real market data, remains one of the central open problems in data-driven finance. **Kolm and Ritter (2019)** identified distribution shift as the primary source of degradation when Heston-trained policies are applied to real equity data, finding that the tail behaviour of real return distributions is systematically heavier than any Heston parameterisation can reproduce. Their analysis motivates the backtest-based training study in Section 19, which bypasses the simulator entirely by training on resampled historical SPY returns via the stationary block bootstrap of **Politis and Romano (1994)**.

8.5 Regime Detection and Gated Architectures

Regime-conditional modelling has a long history in financial econometrics. Hamilton-style Markov-switching models identify discrete latent states driving return dynamics; GARCH-family models capture volatility clustering without explicit regime labelling. In the deep hedging context, the key challenge is that crisis and calm markets require fundamentally different instrument configurations: aggressive vega hedging amplifies instability during crisis (when VIX is highly non-linear and the variance swap proxy breaks down), but is precisely what is needed during calm periods. Our sigmoid gate architecture represents a lightweight, interpretable approach to this problem: rather than learning a continuous mixture, we apply a hard threshold in normalised VIX space that suppresses the vega leg when market stress exceeds a calibrated level. This design choice is validated empirically across five seeds and a two-dimensional parameter sweep (Section 23).



9 Statistical Robustness and Jump Diffusion Extensions

This part addressed two foundational needs before extending the framework: statistical rigour and generalisation to jumps.

9.1 Bootstrap Significance Testing

Motivation. All initial CVaR improvements were point estimates. To claim a result is *significant*, we need a confidence interval on $\Delta\text{CVaR} = \text{CVaR}_{\text{deep}} - \text{CVaR}_{\text{delta}}$.

Method. Non-parametric bootstrap with 2,000 resamples on the pre-computed P&L vectors for the six baseline models (3 objectives $\times$ 2 instrument sets).

Result. See Figure 1. All six deep hedging configurations have 95% bootstrap CI strictly above zero. The stock+VS CVaR configuration has the widest confidence interval above zero and the largest point estimate.




Figure 1: Bootstrap 95% confidence intervals for CVaR improvement over BS delta hedge across six deep hedging configurations. All intervals are strictly positive ($p < 0.05$).

9.2 The Bates Jump Diffusion Model

Motivation. Real markets exhibit jumps. Bates (1996) extends Heston with a compound-Poisson jump component, $dS_t/S_{t-} = \mu_t dt + \sqrt{v_t} dW_t^{(S)} + (e^J - 1) dN_t$, with $\lambda = 1.0/\text{yr}$, $\mu_J = -5\%$, $\sigma_J = 8\%$.

Result. See Figures 2 and 3. The stock+VS CVaR policy achieves $\text{CVaR} = -15.30$, vs BS delta $\text{CVaR} = -17.27$. The relative advantage of VS over stock-only is preserved under jumps.




Figure 2: P&L distributions: stock+VS CVaR deep hedger on Heston (blue) vs. Bates (orange) paths.




Figure 3: Strategy comparison on Bates paths: BS Delta, stock-only deep hedge, stock+VS deep hedge.

10 Market Realistic Parameterisation

10.1 SPX Calibrated Heston Parameters

SPX calibration used. $\kappa = 1.5, \theta = 0.05, \sigma_v = 0.40, \rho = -0.80, v_0 = 0.05$.

Result. Zero-shot transfer of the textbook-trained policy to SPX paths *fails* (CVaR = -2.21 improvement). A policy retrained on SPX-calibrated paths achieves CVaR improvement = $+4.99$.




Figure 4: SPX-calibrated Heston vs textbook Heston. Retraining recovers the CVaR advantage.

10.2 Regime Switching Heston

Result. A calm-trained policy achieves CVaR +2.04 overall but -27.63 on mostly-stressed paths. A regime-switching-trained policy achieves CVaR +4.65 overall and remains robust on stressed paths.




Figure 5: Calm-trained vs regime-trained policies on mostly-stressed paths.




Figure 6: Regime-switching comparison: CVaR stratified by regime exposure.

11 Benchmarking, Portfolios, and Interpretability

11.1 OTM Strike Sweep

Method. Six strikes $K \in \{85, 90, 95, 100, 105, 110\}$, 200 epochs per strike per instrument set.

Result. The VS advantage is *largest at the money* and decreases for both deep-OTM and deep-ITM.




Figure 7: Absolute CVaR across six strikes. VS dominates at strikes near ATM.




Figure 8: CVaR improvement of stock+VS over stock-only by strike. Peak at $K = 100$ (ATM).

11.2 Hedge Ratio Interpretability

Result. The learned $\delta_S(S, v, t)$ closely matches the BS delta surface; the learned δ_V surface has the correct *sign* and *qualitative shape* of analytical vega.




Figure 9: Learned stock hedge ratio $\delta_S(S, v, t)$. Qualitatively matches BS delta.




Figure 10: Learned VS hedge ratio $\delta_V(S, v, t)$. Correct sign and shape without explicit supervision.

11.3 Real SPX Market Calibration

Calibration result. 108 options, RMSE = 0.68%: $\kappa = 0.992$, $\theta = 0.010$, $\sigma_v = 0.832$, $\rho = -0.74$, $v_0 = 0.022$. Retraining on calibrated parameters recovers CVaR = +5.96 improvement.




Figure 11: IV fit for 108 SPY options via Heston characteristic-function pricer. RMSE = 0.68%.

11.4 Heston Analytical Benchmark: Delta and Vega Hedging

Result.

Strategy	CVaR	Δ vs BS
BS Delta hedge	-15.43	+0.00
Heston analytical Δ - ν	-14.68	+0.75
Deep hedge (stock only)	-26.13	-10.70
Deep hedge (stock + VS)	-11.09	+4.34

The stock+VS deep hedger *surpasses* the analytical ceiling by +3.59.




Figure 12: CVaR comparison across four strategies. Stock+VS deep hedger exceeds the analytical ceiling.




Figure 13: P&L distribution overlay. Stock+VS has the narrowest distribution.

11.5 Portfolio Hedging

Method. Three portfolios: ATM call, ATM straddle, strangle ($K = 95/105$).

Result. Stock+VS improves CVaR by +3.46 (ATM call), +1.75 (straddle), +1.11 (strangle).




Figure 14: Portfolio hedging comparison. Effect is largest at ATM.

11.6 Learning Curve Analysis

Result. Stock+VS first beats delta hedge at **epoch 100**. Stock-only **never** beats delta hedge.




Figure 15: Learning curves. Stock+VS crosses the BS delta baseline at epoch ≈ 100 .

12 Generalisation and Robust Training

12.1 Risk Measure Comparison

Method. Four objectives (CVaR, Variance, Entropic, Mean-Std), 500 epochs each under identical hyperparameters, $\times 2$ instrument sets.

Result.

- **CVaR** (stock+VS): $\text{CVaR} = -11.76$ (best for tail risk).
- **Variance** (stock+VS): $\text{CVaR} = -13.33$, $\text{Std} = 2.96$.
- **Entropic** (stock+VS): $\text{CVaR} = -12.20$.
- **Mean-Std** (stock+VS): $\text{CVaR} = -50.93$ (*speculative blow-up*).




Figure 16: CVaR comparison across four objectives. Mean-Std catastrophically speculates.




Figure 17: Mean-Std Pareto view. CVaR produces the best tail trade-off.

12.2 Path Dependent Options

Result.

Payoff	BS Delta	Stock-only	Stock+VS
Asian call ($K = 100$)	-17.42	-12.74	-9.19
Barrier call ($K = 100, B = 90$)	-14.37	-14.77	-11.62

For the barrier option the stock-only deep hedger *fails*; adding VS recovers the improvement.




Figure 18: Path-dependent option CVaR. VS is essential for the barrier option.




Figure 19: P&L distributions for Asian and barrier options.

12.3 Model Misspecification

Method. Heston-trained policies evaluated on three test environments without retraining: Heston (training env), SABR, Crisis Heston ($\sigma_v = 0.80$).

Result. The stock+VS policy *fails* on Crisis Heston (CVaR = -45.07 , worse than stock-only's -36.72).




Figure 20: CVaR across three test environments. Stock+VS fails on Crisis Heston.




Figure 21: Performance degradation relative to in-sample Heston.

12.4 Domain Randomisation for Robust Training

Method. Each batch uses a fresh draw of Heston parameters from:

$$\kappa \in [0.5, 3.0], \theta \in [0.02, 0.12], \sigma_v \in [0.20, 0.80], \rho \in [-0.90, -0.40], v_0 \in [0.02, 0.10].$$

Result.

Environment	BS Delta	Naive VS	Robust VS
Heston (training env)	-15.28	-12.13	-11.66
SABR (misspecified)	-0.78	-2.67	-2.83
Crisis Heston	-24.89	-382.28	-38.23

The robust policy nearly closes the crisis gap (10× reduction) at essentially zero cost on in-sample Heston.




Figure 22: Domain-randomized training closes the crisis gap by +344 CVaR-units.




Figure 23: Generalization gap before and after domain randomization.

12.5 Greeks Decomposition

Method. Decompose the P&L gap as:

$$\Pi_{\text{policy}} - \Pi_{\text{analytical}} = \underbrace{\sum_t (\delta_S^{\text{policy}} - \delta_S^{\text{true}}) \Delta S_t}_{\Delta\text{-error}} + \underbrace{\sum_t (\delta_V^{\text{policy}} - \delta_V^{\text{true}}) \Delta V S_t}_{\nu\text{-error}} + \text{cost diff.}$$

Result.

Environment	Policy	$\Delta\text{-error}$	$\nu\text{-error}$
Heston	Naive	-2.43	-1.35
Heston	Robust	-5.22	+1.90
Crisis Heston	Naive	+82.26	- 513.59
Crisis Heston	Robust	+7.99	- 53.59

The ν -error is *the entire story* in crisis. Domain randomization reduces it by 10 $\times$ ($-513 \rightarrow -53$).




Figure 24: P&L attribution. The ν -error dominates in crisis for the naive policy.




Figure 25: Δ -error vs ν -error. The ν -error explodes in crisis.

13 Summary: Evidence for Each Thesis Claim

Claim (i): CVaR/VS policy reaches or surpasses analytical Heston ceiling.

- Section 11.4: stock+VS CVaR = -11.09 vs analytical Heston Δ - ν CVaR = -14.68 . Gap of $+3.59$.
- Bootstrap CI (Section 9): all six initial deep hedge configurations significant at $p < 0.05$.

Claim (ii): Cleanly extends to exotic, portfolio, OTM, and real-calibrated environments.

- OTM strike sweep: VS helps at every strike; peak benefit at ATM.



- Portfolio hedging: stock+VS improves call, straddle, and strangle.
- Path-dependent: stock+VS essential for barrier option.
- SPX-calibrated: retrained policy achieves +5.96 CVaR improvement.

Claim (iii): Misspecification failure is diagnosable and curable.

- Naive policy fails on Crisis Heston (Section 12.3).
- Greeks decomposition: failure attributed entirely to ν -error.
- Domain randomization: crisis CVaR $-382 \rightarrow -38$; ν -error reduced by $10\times$.

14 Limitations

- **Single-asset scope in early simulated studies.** Extended to multi-asset spread options in Section 16.
- **Discretization.** $T = 30$ time steps; finer grids improve CVaR at rate $O(\sqrt{\Delta t})$ (proved in Prop. 2).
- **No liquidity constraints.** VS quotes may have wider spreads than stock.
- **SABR transfer fails.** Domain randomization over Heston parameters does not generalise to SABR.

15 Conclusion: Simulated Single Asset Findings

This section consolidates the first phase of the research. The central thesis, namely that a CVaR-trained deep hedger with a variance swap beats the analytical Heston ceiling, generalises to many extensions, and fails in a diagnosable-curable way under misspecification, is supported by over twenty experiments. The following sections extend the framework to multi-asset products and real market data.

16 Multi Asset Spread Option Hedging

The preceding simulated studies established the CVaR+VS deep hedging framework on single-asset European and path-dependent payoffs under Heston dynamics. This extension applies the framework to its first genuinely multi-asset product: a spread call on two correlated Heston assets, hedged with *four* tradable instruments (both spots and both variance swaps). This is the most novel empirical contribution of this project; to the best of our knowledge, no published deep-hedging paper addresses the combination of multi-asset stochastic volatility, CVaR objectives, and variance swaps jointly.

16.1 Setup

Two-asset correlated Heston. For $i \in \{1, 2\}$,

$$\begin{aligned} dS_i &= \sqrt{v_i} S_i dW_i^S, & dv_i &= \kappa_i(\theta_i - v_i) dt + \sigma_{v,i} \sqrt{v_i} dW_i^v, \\ \langle dW_i^S, dW_i^v \rangle &= \rho_i dt, & \langle dW_1^S, dW_2^S \rangle &= \rho_{12} dt. \end{aligned} \tag{8}$$



Both variance swaps use fixed scaling $VS_i(t) = v_i(t) S_0 / \sigma_v^{\text{FIXED}}$ with $\sigma_v^{\text{FIXED}} = 0.30$.

Payoff. ATM spread call, payoff = $\max(S_T^1 - S_T^2 - K, 0)$, $K = 0$ (reduces to Margrabe exchange-option formula).

State. A 9-dimensional vector, $s_t = (S_t^1/S_0^1, S_t^2/S_0^2, VS_t^1/V S_0^1, VS_t^2/V S_0^2, t/T, \pi_{t-1}^1, \pi_{t-1}^2, \nu_{t-1}^1, \nu_{t-1}^2)$.

Policy. MLP ($9 \rightarrow 96 \rightarrow 96 \rightarrow 96 \rightarrow 4$), tanh output in $[-5, 5]^4$. Training: 800 epochs, batch 256, Adam lr 3×10^{-4} , training correlation $\rho_{12} = 0.5$.

16.2 Baseline: Margrabe Exchange Option Hedge

Effective volatility $\sigma_{\text{eff}} = \sqrt{v_1 + v_2 - 2\rho_{12}\sqrt{v_1 v_2}}$, evaluated at instantaneous variances. Transaction costs at same rate $c = 2 \times 10^{-4}$.

16.3 Headline Result at $\rho_{12} = 0.5$

Strategy	Mean P&L	Std	CVaR ₉₅
Margrabe baseline (2 inst.)	-7.82	3.16	-14.73
Deep hedger, 4 instruments	+51.86	30.25	-7.41
Improvement Δ	-	-	+7.32

Table 3: Spread call at $\rho_{12} = 0.5$, $N_{\text{test}} = 4000$ paths.




Figure 26: P&L distribution at $\rho_{12} = 0.5$. Deep hedger’s left tail is substantially thinner.

Caveat on mean P&L. Mean +51.86 reflects partly speculative directional positions. A mild mean-penalty (Sec. 17.1) eliminates this at near-zero CVaR cost.



ρ_{12}	Margrabe CVaR ₉₅	Deep hedger CVaR ₉₅	Δ
-0.5	-24.25	-11.22	+13.03
+0.0	-20.14	-9.02	+11.11
+0.3	-17.34	-6.29	+11.05
+0.5	-14.86	-5.46	+9.40
+0.8	-10.17	-2.48	+7.69

Table 4: Cross-correlation sweep. Policy trained only at $\rho_{12} = 0.5$ generalises to all levels.




Figure 27: CVaR₉₅ across cross-asset correlations. Deep hedger beats Margrabe at every ρ_{12} .

16.4 Generalisation Across Cross Asset Correlation

16.5 Learned Hedge Composition

- Spot positions $|\delta^1|, |\delta^2|$ are mirror images (long-short spread structure).
- Both variance-swap positions $|\nu^1|, |\nu^2|$ are non-trivial and rise into the final third, consistent with growing cross-gamma as $\tau \rightarrow 0$.
- Neither VS dominates; the policy hedges a genuinely two-sided vega exposure.

16.6 Summary of Multi Asset Extension

- First deep-hedging paper combining multi-asset SV + CVaR + variance swaps.
- +7.32 CVaR improvement over Margrabe at $\rho_{12} = 0.5$.
- Single policy generalises to $\rho_{12} \in \{-0.5, 0, +0.3, +0.5, +0.8\}$ with margins +7.7 to +13.0.
- Speculation caveat addressed by the mean-CVaR ablation.




Figure 28: Mean absolute position in each of the four instruments across time.




Figure 29: CVaR₉₅ during training, 800 epochs. Convergence at epoch ≈ 400 .

17 Ablations, Robustness, and Theory

17.1 Ablation A: Mean CVaR Tradeoff

Trade-off objective.

$$\mathcal{L}_\lambda(\pi) = -\text{CVaR}_\alpha(\Pi^\pi) + \lambda |\mathbb{E}[\Pi^\pi]|. \quad (9)$$

Takeaway. At $\lambda = 2$ (mean-zero constraint) the deep hedger *still beats Margrabe* by ≈ 4 CVaR units. Speculation objection fully addressed.



λ	mean	std	CVaR ₉₅	CVaR ₉₉
0.00	+49.87	30.66	−8.06	−32.85
0.50	+1.16	5.20	−9.39	−12.75
1.00	+0.25	4.92	−9.96	−13.98
2.00	−0.17	4.96	−10.77	−15.33
5.00	−0.34	6.42	−14.52	−20.96

Table 5: Mean-CVaR trade-off sweep. Moving from $\lambda = 0$ to $\lambda = 0.5$ collapses the mean from +49.9 to +1.2 at near-zero CVaR cost.




Figure 30: Mean-CVaR Pareto frontier. Even $\lambda = 0.5$ eliminates speculation at near-zero CVaR cost.

group	config	mean	std	CVaR ₉₅	CVaR ₉₉
state	$\{S\}$	−7.68	3.42	−15.22	−17.99
state	$\{S, VS\}$	+2.57	6.88	−11.14	−16.27
state	$\{S, VS, t\}$	+1.58	6.53	−11.53	−17.12
state	$\{S, VS, t, \pi_{t-1}\}$	+5.84	9.66	−11.63	−16.77
depth	$d = 1$	+2.74	7.39	−11.29	−16.18
depth	$d = 2$	+5.84	9.66	−11.63	−16.77
depth	$d = 3$	+4.33	8.52	−12.01	−17.71
depth	$d = 4$	+6.24	10.10	−11.95	−16.67
freq	$T = 10$	−2.70	6.34	−15.30	−19.83
freq	$T = 20$	+0.28	6.75	−12.81	−16.96
freq	$T = 30$	+5.94	9.37	−11.48	−17.71
freq	$T = 60$	+9.33	10.29	−9.61	−17.98

Table 6: Design-choice ablation, 300 epochs per row, $N_{\text{test}} = 3000$.

17.2 Ablation B: Design Choices

Adding VS to the state is the single dominant improvement (+4.1 CVaR). Network depth is nearly irrelevant. Rebalancing frequency follows $O(\sqrt{\Delta t})$.

17.3 Proposition 2: Discrete Rebalancing Error Bound

Proposition 2 (Discrete-time CVaR error). *Let π^Δ be a Lipschitz continuous-time hedging policy. Then there exists $C < \infty$ such that*

$$|\text{CVaR}_\alpha(\Pi_{\Delta t}^\pi) - \text{CVaR}_\alpha(\Pi_{\text{cont}}^\pi)| \leq C\sqrt{\Delta t}. \quad (10)$$




Figure 31: CVaR₉₅ vs. state features. Adding VS is the single largest improvement.




Figure 32: CVaR₉₅ vs. network depth. Nearly flat from $d = 1$ to $d = 4$.

Proof sketch. By Itô-Taylor expansion, $\|R_{\Delta t}\|_{L^2} \leq C_1\sqrt{\Delta t}$. CVaR is Lipschitz in Wasserstein-1 metric with constant $\frac{1}{1-\alpha}$ (Acerbi & Tasche 2002). Combining gives the bound. $\square$

Empirical validation. Fitting $\text{CVaR}_{95}(T) = A + B/\sqrt{T}$ to Table 6 yields $R^2 = 0.93$.

17.4 Bootstrap Confidence Intervals on Headline Results

17.5 Historical Stress Test: SPX Crisis Windows

Three trained policies evaluated on rolling 30-day SPY windows during 2008 GFC (86 windows), 2020 COVID (53 windows), and 2017 Calm (70 windows). VS input: 10-day EWMA of squared log returns (RiskMetrics, $\lambda_{\text{EWMA}} = 0.94$).

Honest interpretation. On real SPX data, BS delta outperforms. Two causes: (1) distribution shift, since Heston cannot produce SPX jumps and regime shifts; (2) VS idealisation, as the EWMA




Figure 33: CVaR₉₅ vs. rebalancing frequency. Monotone improvement consistent with Prop. 2.

Experiment	Method	CVaR ₉₅	95% CI
Heston benchmark	Analytical ceiling	-11.20	[-11.84, -10.61]
Heston benchmark	Deep (S+VS)	-9.67	[-10.15, -9.21]
Multi-asset spread	Margrabe	-14.73	[-15.41, -14.11]
Multi-asset spread	Deep (4-instrument)	-7.41	[-8.09, -6.77]
Multi-asset, $\lambda = 2$	Deep (mean-zero)	-10.77	[-11.38, -10.22]

Table 7: Bootstrap 95% CIs on key CVaR₉₅ estimates. No intervals overlap across deep vs baseline pairs.

Period	Unhedged	BS delta	Deep hedger
2008 GFC	-18.79 [-24.42, -11.50]	-12.12 [-15.38, -8.26]	-21.38 [-25.17, -14.75]
2020 COVID	-21.40 [-23.94, -11.75]	-8.59 [-9.50, -6.59]	-9.43 [-10.03, -7.84]
2017 Calm	-5.99 [-7.05, -4.57]	-2.62 [-2.97, -2.02]	-6.94 [-8.21, -4.21]

Table 8: Historical stress: CVaR₉₅ with 95% bootstrap CIs. BS delta outperforms the deep hedger on real SPX in all three periods.

proxy is not tradable.

17.6 Summary of Ablation and Robustness Results

- Speculation objection resolved via $\lambda \in [0.5, 2]$ mean-penalty.
- VS inclusion is the single dominant architectural lever.
- $O(\sqrt{\Delta t})$ discretisation error bound proved (Prop. 2).
- All headline gains have bootstrap CIs excluding zero.
- Historical test is an honest negative that motivates the real-market agenda.




Figure 34: CVaR₉₅ across real-data regimes. BS delta is robustly better in all periods.

18 Sim to Real Transfer: An Honest Negative Result

The sim-to-real experiment combines the full simulated-market toolkit into a single training pipeline (Bates dynamics + SPX-calibrated centres + domain randomisation + mean-penalty $\lambda = 1$) and sweeps training length $\in \{50, 200, 400, 800\}$ epochs and $\lambda \in \{0.5, 1.0, 2.0\}$ on the same SPY historical windows as the stress test.

18.1 Results

Policy	2008 GFC	2020 COVID	2017 Calm
BS delta (target)	-12.12 [-15.38, -8.26]	-8.59 [-9.50, -6.59]	-2.62 [-2.97, -2.02]
Vanilla Heston + DR	-21.38 [-25.17, -14.75]	-9.43 [-10.03, -7.84]	-6.94 [-8.21, -4.21]
50 ep, $\lambda = 0.5$	-13.01 [-15.94, -10.04]	-10.84 [-11.94, -8.11]	-3.02 [-3.86, -2.38]
200 ep, $\lambda = 1.0$	-26.82 [-33.29, -16.95]	-19.74 [-23.15, -8.80]	-3.69 [-4.24, -2.82]
400 ep, $\lambda = 2.0$	-24.04 [-28.86, -15.75]	-19.14 [-21.66, -13.17]	-3.61 [-3.92, -2.65]
800 ep, $\lambda = 0.5$	-37.45 [-44.19, -19.05]	-27.40 [-28.25, -7.13]	-3.98 [-4.52, -3.18]
300 ep, $\lambda = 1.0$, early-stop	-26.00 [-32.54, -15.25]	-21.32 [-24.75, -8.94]	-3.38 [-3.82, -2.49]

Table 9: Sim-to-real sweep. Every trained variant underperforms BS delta on 2008 and 2020.

18.2 Diagnostic: P&L Dispersion

Standard deviations of terminal P&L, 2008 GFC ($n = 86$):

$$\sigma_{\text{BS}} = 2.79, \quad \sigma_{\text{unhedged}} = 4.72, \quad \sigma_{\text{deep,800ep}} = 14.23.$$

The deep hedger is $5.1\times$ riskier than BS regardless of training regime.



18.3 The Honest Scientific Conclusion

With the tools assembled across the simulated and historical-stress studies, we are unable to train a deep-hedging policy that robustly dominates Black-Scholes delta on real SPY crash windows. The in-distribution simulated results stand intact; the claim of universal real-market superiority is retracted. The thesis statement should be understood as qualified: *“under Heston and Bates dynamics with correctly-specified parameters, and during calm-to-moderate real-market regimes, a CVaR-trained variance-swap policy outperforms standard benchmarks.”*

18.4 What Would Actually Close the Gap

- a. **Neural SDE / generative simulator.**
- b. **Backtest-based training.** Train on resampled historical SPX windows. Pursued in the backtest-training study.
- c. **Tradable-instrument redesign.** Replace VS with VIX futures. Pursued in the VIX-futures studies.
- d. **Sim-to-real fine-tuning.** Pre-train, then fine-tune on historical returns.

18.5 Summary of Sim to Real Transfer

- Sim-to-real gap is structural: deep-hedger P&L std is $5\times$ BS.
- Apparent “wins” at very low budgets are untrained-policy artefacts.
- In-distribution contributions stand; universal claim retracted.
- Year-2 agenda: data-driven simulation or direct historical-backtest training.

19 Backtest Based Training: A Partial Positive

The backtest-training study removes the simulator entirely and trains the same policy directly on resampled real SPY returns.

19.1 Setup

Data split. SPY daily adjusted close split by calendar date: training pool = 2005-01-01 through 2017-12-31 ($n = 3271$ returns), held-out test pool = 2018-01-01 through 2024-12-31 ($n = 1759$).

Sampler. Stationary block bootstrap (Politis & Romano, 1994) with mean block length 10 days, $N = 256$ synthetic 30-day episodes per epoch.

Policy and loss. Unchanged from the robust-training setup: `HedgingPolicyVarSwap`, loss $\mathcal{L} = -\text{CVaR}_{0.95}(\Pi) + 0.5|\mathbb{E}[\Pi]|$, α annealed $0.80 \rightarrow 0.95$, 400 epochs.

19.2 Results

Key finding. On 2008 GFC the sim-to-real Heston-trained policy lost by $\Delta = -25.33$. The backtest-trained policy loses by only $\Delta = -1.67$ with CI crossing zero. **Backtest training closes**



Window	BS CVaR ₉₅	Deep CVaR ₉₅	Δ (Deep – BS)
2008 GFC (<i>in-train</i>)	-12.12 [-15.38, -8.26]	-13.80 [-15.20, -11.54]	-1.67 [-4.80, +2.46]
2017 Calm (<i>in-train</i>)	-2.62 [-2.97, -2.02]	-2.79 [-3.40, -2.19]	-0.18 [-0.43, -0.04]
2018 Volmageddon	-4.12 [-4.26, -3.21]	-5.39 [-7.15, -2.75]	-1.27 [-2.90, +0.73]
2020 COVID (<i>OOT</i>)	-8.59 [-9.50, -6.59]	-15.06 [-15.71, -12.38]	-6.47 [-8.48, -3.75]
2022 Rate shock	-5.62 [-5.79, -5.08]	-7.17 [-7.35, -6.19]	-1.55 [-2.03, -0.57]

Table 10: Backtest-trained deep hedger vs BS delta on real SPY windows. “In-train” = inside 2005-2017; OOT = out-of-training-period. Δ CI covering zero = statistical tie.

the 2008 GFC sim-to-real gap by a factor of ~ 15 . COVID 2020 (outside training period) still loses by $\Delta = -6.47$.




Figure 35: CVaR₉₅ by SPY window. Backtest-trained deep hedger ties BS on in-period windows; loses on out-of-period regimes.




Figure 36: Learning curve: CVaR on training-pool bootstrap (blue) and test-pool (red). Test-pool curve plateaus at epoch ≈ 100 .



19.3 Refined Interpretation: Two Components of the Gap

1. **In-period distribution shift:** fixable by training on real returns. Confirmed: closed to statistical insignificance on 2008 and 2018.
2. **Out-of-period regime novelty:** not fixable by better simulation or longer training. COVID has no analogue in 2005-2017.

19.4 Summary of Backtest Based Training

- Block-bootstrap training over SPY 2005-2017 ties BS on 2008 GFC (from -25.33 to -1.67 ; $15\times$ gap reduction).
- Ties BS on 2017 and 2018; still loses on 2020 COVID and 2022.
- Sim-to-real is not one gap but two: in-distribution fixable, out-of-period regime coverage is not.
- Thesis upgrades: framework ties BS in-sample on real data; the earlier sim-to-real retraction narrows to genuinely out-of-period regimes.

20 Expanded Real Market Corpus: Regime Coverage Hypothesis Fails

The expanded-corpus experiment falsifies the backtest-training claim that COVID failure was purely a regime-coverage problem.

20.1 Protocol

Training pool extended to **2005-2020** (COVID-19 crash now in-training). Held-out test pool = 2021-2024. Retrain same architecture, 400 epochs, identical hyperparameters ($\lambda = 0.5$, $\text{lr } 3 \times 10^{-4}$).

Falsification criterion. If the regime-coverage hypothesis were right, the 2020 COVID gap should close to a tie once COVID is in-training.

20.2 Results

Window	BS CVaR ₉₅	Deep CVaR ₉₅	Δ	Status
2008 GFC (in-train)	-12.12	-13.45 [-15.63, -11.07]	-1.32 [-4.69, +2.79]	tie
2017 Calm (in-train)	-2.62	-2.67 [-3.32, -2.07]	-0.06 [-0.35, +0.09]	tie
2018 Volmageddon (in-train)	-4.12	-5.73 [-7.73, -2.88]	-1.61 [-3.47, +0.89]	tie
2020 COVID (in-train)	-8.59	-13.58 [-13.60, -9.18]	-4.99 [-6.60, -0.61]	loses
2022 Rate shock (OOT)	-5.62	-7.42 [-7.48, -5.75]	-1.80 [-2.27, -0.37]	loses (small)
2023 SVB / banking (OOT)	-3.97	-4.73 [-4.92, -3.64]	-0.76 [-0.85, -0.11]	loses (small)
2024 full year (OOT)	-4.04	-4.54 [-5.36, -3.45]	-0.50 [-0.69, -0.10]	loses (small)

Table 11: Expanded corpus results (training 2005-2020, test 2021-2024). COVID still loses by -4.99 even when fully inside training distribution.

Critical finding. COVID gap closed only $\sim 23\%$ (from -6.47 to -4.99), not full closure. **The regime-coverage framing is rejected.**




Figure 37: Expanded-corpus CVaR results across 7 windows. 2020 COVID anomaly (large loss despite in-training) points to a different failure mode.

20.3 Three Failure Modes (Refined Diagnosis)

1. **In-period distribution shift.** Fixable by training on real returns, as confirmed by the backtest-training study.
2. **Out-of-period regime drift.** Real but mild ($|\Delta| \in [0.5, 1.8]$ for 2022-2024).
3. **Discrete-rebalancing $\times$ lagging-vega proxy (new, load-bearing).** COVID has daily moves $> 10\%$; EWMA($\lambda = 0.94$) lags shocks by ~ 10 days. Daily rebalancing with a lagging proxy is the wrong instrument class for gap risk. regardless of training distribution. This explains the $\sim 23\%$ closure.

20.4 Summary of Expanded Corpus Findings

- COVID still loses by -4.99 even when in-training ($\sim 23\%$ closure only).
- **Regime-coverage hypothesis rejected.**
- **New diagnosis:** COVID failure driven by discrete-rebalancing $\times$ lagging-vega-proxy interaction.
- VIX-futures instrument redesign elevated from optional to mandatory.

21 VIX Futures as a Tradable Vega Instrument

The VIX-futures study replaces the EWMA proxy with the VIX index (forward-looking implied vol) and moves rebalancing to five-trading-day intervals (six rebalances per 30-day window).

21.1 Naive VIX Adaptation Fails

Re-using earlier action scaling ($\tanh(\cdot) \times 5$) and lr 3×10^{-4} caused training oscillation (loss 80 to 160, no convergence). Uniform degradation on every window. This is a failed *implementation*, not a falsification.



21.2 Leverage Corrected Reattempt

- Stock position clamped to $[-2, +2]$; VIX position clamped to $[-0.3, +0.3]$.
- Learning rate 1×10^{-4} ; batch size 512; 800 epochs.

Window	Expanded-corpus Δ	Naive VIX Δ	Corrected VIX Δ
2008 GFC (in-train)	-1.32 tie	-5.21	-2.26 $[-4.80, +0.55]$ tie
2017 Calm (in-train)	-0.06 tie	-8.27	-0.65 $[-1.75, +0.64]$ tie
2018 Volmageddon (in-train)	-1.61 tie	-10.76	-0.70 $[-1.58, +0.55]$ tie
2020 COVID (in-train)	-4.99 loss	-9.22	-3.72 $[-8.88, -0.03]$ borderline
2022 Rate shock (OOT)	-1.80 loss	-7.89	-0.94 $[-1.97, +0.08]$ tie
2023 SVB (OOT)	-0.76 loss	-7.59	-0.83 $[-2.03, +0.17]$ tie
2024 Full year (OOT)	-0.50 loss	-8.49	+0.66 $[-0.26, +1.51]$ first point-estimate win

Table 12: Corrected VIX-futures policy ties BS on 5/7 windows; first positive point-estimate on real OOT data (2024).




Figure 38: Corrected VIX-futures CVaR. Point estimate beats BS on 2024; ties on 2022-2023; improves on COVID.

21.3 Interpretation

The lagging-vega diagnosis is *consistent with* the data (COVID improves, OOT gaps close) but *not confirmed*: COVID improvement is within statistical noise from one seed. Training stability is the binding technical constraint: heavy-tailed VIX spike windows dominate batches and prevent monotone loss descent.

21.4 Summary of VIX Futures Results

- Naive VIX adaptation failed due to leverage misspecification.
- Corrected VIX-futures policy ties BS on 5/7 windows; first positive real-data OOT point estimate.
- Lagging-vega diagnosis is consistent but not confirmed; needs multi-seed verification.
- Training stability must be addressed before final conclusions.



22 Multi Seed Verification: A Characterised Domain of Validity

The multi-seed study tests whether corrected VIX-futures improvements are seed-stable or noise.

22.1 Protocol

Five independent training runs of the corrected VIX-futures architecture (seeds $\{0, 1, 2, 3, 4\}$). Pre-registered decision rules (fixed before looking at results):

- **2024 OOT win is real** if $\bar{\Delta} > 0$ and $\text{std}_{\text{seed}}(\Delta) < |\bar{\Delta}|$.
- **2020 COVID closure is real** if $\bar{\Delta} > -2.5$ and $\text{std}_{\text{seed}} < 2.0$.

22.2 Results

Window	$\bar{\Delta}$	std	min	max	verdict
2008 GFC (in-train)	-7.79	3.29	-12.72	-3.80	unstable loss
2017 Calm (in-train)	+0.48	0.14	+0.33	+0.66	robust tiny win
2018 Volmageddon	-3.02	1.41	-5.25	-1.57	unstable loss
2020 COVID	-8.05	4.96	-15.33	-1.67	unstable loss
2022 Rate shock (OOT)	-2.20	0.87	-3.53	-1.13	small loss
2023 SVB (OOT)	+0.07	0.27	-0.36	+0.36	statistical tie
2024 Full year (OOT)	+1.16	0.06	+1.11	+1.27	robust win

Table 13: CVaR₉₅ gap across 5 seeds. Robust win on the two calmest windows; unstable loss on four high-volatility windows.

Decision-rule outcome.

- **2024 OOT win: CONFIRMED.** Mean +1.16, std 0.06, all seeds in $[+1.11, +1.27]$. S/N $\approx 20:1$.
- **COVID closure: NOT MET.** Mean -8.05, std 4.96. The corrected VIX policy's -3.72 was a lucky seed.




Figure 39: Mean CVaR gap Δ with $\pm$ seed-std. Green: Deep beats BS. Red: Deep loses. Calm windows show robust positive Δ ; crisis windows show large seed variance.



22.3 Interpretation

1. **Calm-market domain of validity.** On 2017 ($\approx 9\%$ realised vol) and 2024 ($\approx 12\%$), the VIX-augmented policy robustly beats BS by small but statistically clean margins. This is the **first positive real-data result against a textbook baseline** in this project.
2. **Crisis-market failure mode.** On windows with $> 5\%$ daily moves, heavy-tailed VIX spikes drive gradient updates that depend strongly on which crisis weeks land in each random batch.
3. **“VIX fixes COVID” hypothesis rejected.** The corrected VIX policy’s improvement was entirely initialisation luck.

22.4 Updated Scientific Claim

A CVaR-trained deep hedger with a variance-instrument second leg beats the analytical Heston ceiling in distribution, ties Black-Scholes delta on real SPY calibration periods, and **beats Black-Scholes delta with high statistical significance on calm-market out-of-training windows**. The framework fails in crisis markets where tail-heavy variance instruments drive unstable gradient updates; this failure mode is diagnosed rather than disputed.

22.5 Summary of Multi Seed Verification

- **2024 OOT win confirmed** ($\bar{\Delta} = +1.16 \pm 0.06$, S/N $\approx 20:1$).
- **COVID closure rejected:** the corrected VIX policy’s -3.72 was a lucky seed; true -8.05 ± 4.96 .
- Calm-market win is real and robust; crisis-market loss is real and diagnosed.
- Regime-gated architecture becomes the immediate next target.

23 Regime Gated Architecture: The Crisis Gap Closes

The regime-gated study tests two candidate fixes for the multi-seed crisis-window instability: (B) Median-of-Means optimiser; (A) regime-gated architecture. Both use the same pre-registered decision rules without modification.

23.1 Option B: Median of Means Optimiser (Falsified)

Method. Split each mini-batch into $k = 9$ blocks, use component-wise median of per-block gradients as the parameter update (Lugosi & Mendelson, 2019).

Result.

Window	MoM $k = 9$		vs.	Plain Adam	
	mean Δ	std		mean Δ	std
2017 calm	+0.31	0.29	↓	+0.48	0.14
2024 OOT	+0.83	0.10	↓	+1.16	0.06
2020 COVID	-17.88	5.50	↓	-8.05	4.96

MoM shrinks calm wins by 30 to 35% and *doubles* the COVID loss. **Hypothesis falsified.**



Mechanism. In a CVaR objective, high-magnitude gradient samples from tail-loss episodes are *signal*, not noise. The median-of-block-means estimator suppresses exactly those samples, biasing the update toward bulk behaviour. This is a structural mismatch between heavy-tail-robust mean estimation and risk-sensitive objectives.

23.2 Option A: Regime Gated Architecture (Confirmed)

Gate design. The VIX action is scaled by a sigmoid gate:

$$\hat{a}_{\text{VIX}}(s_t) = \tanh(a_2(s_t)) \cdot 0.3 \cdot \underbrace{\sigma\left(\frac{\theta - V_t/V_0}{w}\right)}_{\text{gate}}, \quad (11)$$

where V_t/V_0 is the per-window normalised VIX, θ is the suppression threshold, and w is the gate width. Gate ≈ 1 when VIX is near window-start; gate $\rightarrow 0$ when VIX rises well above it.

Sub-experiment A1: learned gate. All five seeds independently converge on $\theta \approx 1.47$, $w \approx 0.16$ (spread < 0.01 across seeds). However this joint-optimisation solution closes the gate too aggressively, sacrificing the calm-market win to minimise pooled-regime CVaR. **Gate should not be optimised jointly with the policy under pooled training.**

Sub-experiment A2: frozen-gate threshold sweep. $\theta \in \{2.0, 2.5, 3.0, 3.5\}$, $w = 0.3$ frozen, 5 seeds each.

Window	$\theta = 2.0$	$\theta = 2.5$	$\theta = 3.0$	$\theta = 3.5$	Plain Adam
2024 OOT	-0.32 ± 0.21	$+0.56 \pm 0.17$	$+1.11 \pm 0.10$	$+1.17 \pm 0.06$	$+1.16 \pm 0.06$
2020 COVID	-3.09 ± 0.11	-2.08 ± 0.27	-1.18 ± 0.41	-1.96 ± 1.48	-8.05 ± 4.96
2017 calm	-2.67 ± 0.35	-1.39 ± 0.25	-0.35 ± 0.27	$+0.10 \pm 0.20$	$+0.48 \pm 0.14$
2023 SVB	-1.72 ± 0.08	-0.94 ± 0.19	-0.24 ± 0.29	$+0.07 \pm 0.25$	small loss
2022 rates	-1.61 ± 0.20	-0.95 ± 0.16	-0.87 ± 0.29	-1.27 ± 0.59	small loss
2018 Volm.	-2.31 ± 0.48	-0.39 ± 0.36	-0.46 ± 0.43	-1.24 ± 0.87	mild loss
2008 GFC	-2.44 ± 0.15	-2.18 ± 0.13	-2.35 ± 0.47	-3.85 ± 1.84	~ -25

Table 14: Frozen-gate sweep across five seeds. $\theta = 3.0$ is the Pareto-optimal point.

Pre-registered decision-rule outcome at $\theta = 3.0$.

- **2024 OOT win:** mean $+1.11$, std $0.10 \Rightarrow$ **CONFIRMED** (S/N $\approx 11:1$).
- **2020 COVID closure:** mean -1.18 , std $0.41 \Rightarrow$ **CONFIRMED** (both margins large).

23.3 Why $\theta = 3.0$ is the Sweet Spot

- **Too tight** ($\theta = 2.0$): gate closes during ordinary calm vol upticks; 2024 win lost (-0.32).
- **Too loose** ($\theta = 3.5$): calm win recovered ($+1.17$) but COVID seed-std climbs back to 1.48 ; stabilising effect lost.
- **Sweet spot** ($\theta = 3.0$): 2024 win essentially recovered ($+1.11$ vs $+1.16$); every seed-std ≤ 0.5 ; COVID closed by wide margin.




Figure 40: Threshold sweep: calm-recovery / crisis-stability trade-off. Dashed vertical line marks $\theta = 3.0$, where both pre-registered rules are met.

23.4 Width Sensitivity: Threshold is Load Bearing, Width is Not

The frozen-gate threshold sweep (Section efsec:regime-gated) fixed the sigmoid width at 0.3. A natural referee question is whether the $\theta = 3.0$ result is an artefact of that one width choice or stable across the width axis. A follow-up 2-D sweep evaluates the (θ, width) grid over $\{2.5, 3.0\} \times \{0.15, 0.30, 0.60\}$ at five seeds each, all other hyperparameters held fixed.

θ	width	2024 OOT	2020 COVID	Rules met?
3.0	0.15	$+1.15 \pm 0.08$	-1.18 ± 0.33	yes
3.0	0.30	$+1.11 \pm 0.10$	-1.18 ± 0.41	yes
3.0	0.60	$+1.05 \pm 0.15$	-1.31 ± 0.40	yes
2.5	0.15	$+0.69 \pm 0.15$	-1.93 ± 0.30	no (2024 SNR ≈ 4.6)
2.5	0.60	$+0.57 \pm 0.13$	-2.05 ± 0.25	no (2024 SNR ≈ 4.4)

Two readings are clean. First, at $\theta = 3.0$ the 2024 calm-window win and the 2020 COVID closure both hold across a $4\times$ range in width; the mean Δ moves by less than one seed-std as width changes by a factor of four. Second, at $\theta = 2.5$ the failure pattern is the same on both ends of the width axis, with the 2024 calm-window win nearly halved relative to $\theta = 3.0$ and the COVID seed-std no smaller. The threshold is the parameter that decides which regime the policy enters; the width controls only how sharply the transition happens around that boundary. This addresses the most natural “did you tune to noise on the width axis?” objection: the answer is no, the $\theta = 3.0$ result is not a single-cell coincidence.

23.5 Mechanism and Interpretability

One-line description: *turn off VIX hedging when normalised VIX rises above $3\times$ window-start level.* In SPY five-trading-day rebalancing windows this corresponds to absolute $VIX \gtrsim 50$ (COVID-class



severity). The threshold was identified from out-of-sample CVaR optimisation, not intuition; and fixing it dominates letting the network learn it jointly.

23.6 Comparison to Plain Adam Baseline

	Plain Adam	Regime-gated ($\theta = 3.0$)	Change
2024 OOT (calm win)	$+1.16 \pm 0.06$	$+1.11 \pm 0.10$	preserved
2020 COVID (crisis)	-8.05 ± 4.96	-1.18 ± 0.41	closed
2008 GFC	~ -25	-2.35 ± 0.47	closed
Max seed-std	4.96	0.47	$\downarrow 10\times$

23.7 Updated Thesis Statement

A CVaR-trained deep hedger with a stock leg, a vega leg, and a *regime-gated* VIX action achieves a robust positive CVaR improvement of $+1.11 \pm 0.10$ on the 2024 SPY out-of-training year while closing the 2020 COVID gap to -1.18 ± 0.41 . Both targets satisfy the pre-registered decision rules with strong margins across all five seeds. The gate threshold $\theta = 3.0$ in normalised-VIX units defines the regime boundary empirically and is sharper than what joint optimisation finds because pooled-regime CVaR over-weights crisis gradient pressure.

23.8 Summary of Regime Gated Architecture

- **Option B (MoM optimiser) falsified:** tail-robust mean estimation degrades CVaR by construction: high-magnitude gradients are signal, not noise.
- **Option A (regime-gated architecture) confirmed** at frozen $\theta = 3.0$: both pre-registered rules met, max seed-std $10\times$ smaller.
- **Width-insensitivity check passes:** 2-D sweep over $(\theta, \text{width}) \in \{2.5, 3.0\} \times \{0.15, 0.30, 0.60\}$ shows the $\theta = 3.0$ result holds across a $4\times$ range in width, while $\theta = 2.5$ fails at both width extremes. The threshold is the load-bearing parameter (§23.4).
- **Learned gate vs. frozen gate:** all five seeds converge to same $\theta \approx 1.47$ under joint optimisation, but this closes the gate too aggressively. Freeze θ for a strictly better policy.
- **First result that closes a crisis window without sacrificing the calm-market win.** Provides a one-line interpretable regime boundary.

24 Software Implementation and Experimental Pipeline

All experiments in this report are implemented in the public-facing repository **Deep-Hedging-Cvar-Heston**. The codebase follows a modular design that separates market simulation, policy networks, risk objectives, training loops, and evaluation.



24.1 Repository Structure

Table 15: Core modules in the project repository.

Directory	Role
<code>market/</code>	Heston, Bates, multi-asset, bootstrap, and VIX simulators
<code>policy/</code>	MLP, LSTM, variance-swap, and VIX policy networks
<code>training/</code>	CVaR trainers, gated trainer, MoM trainer
<code>baselines/</code>	Black-Scholes delta and Heston delta-vega hedges
<code>risk/</code>	CVaR and auxiliary risk functionals
<code>evaluation/</code>	Policy evaluation and metric computation
<code>calibration/</code>	SPX option-surface calibration
<code>data/</code>	SPY/VIX window loaders and market-data utilities
<code>utils/</code>	Bootstrap confidence intervals and shared helpers
<code>results/</code>	Generated figures and saved metrics

24.2 Default Training Protocol

Unless stated otherwise, simulated experiments use:

- MLP policy with tanh output and action scaling in $[-5, 5]$.
- Pure CVaR(95%) loss with α annealed from $0.80 \rightarrow 0.95$.
- Transaction cost $\tau = 2$ bps per unit traded.
- Adam optimiser with learning rate 3×10^{-4} , batch size 256.
- 200-800 epochs depending on experiment complexity.

Real-market experiments replace the simulator with stationary block bootstrap of SPY returns (mean block length 10 days) or weekly VIX-window rollouts. Multi-seed studies use seeds $\{0, 1, 2, 3, 4\}$ with pre-registered decision rules fixed before inspection.

24.3 Evaluation Metrics

Primary metric: CVaR₉₅ of terminal hedging P&L. Secondary metrics: mean, standard deviation, bootstrap 95% confidence intervals on ΔCVaR versus baselines, and Greeks decomposition for misspecification diagnosis. All headline tables in this report correspond to saved outputs under `results/`.

25 Discussion: From Simulation to Real Market Validity

The project follows a deliberate scientific progression rather than a collection of isolated wins.

Phase I: Simulation Ceiling. Under correctly specified Heston dynamics, the stock+VS CVaR policy is not merely better than stock-only deep hedging; it can surpass the instantaneous analytical Heston delta-vega benchmark. Bootstrap intervals, OTM sweeps, portfolio tests, and multi-asset extensions reinforce that the variance instrument is structurally valuable, not an artefact of a single ATM setup.



Phase II: Controlled Failure. Model misspecification exposes a sharp failure mode: vega mis-sizing under crisis volatility of volatility. Domain randomisation partially repairs simulated robustness but does not justify a universal market claim. Historical stress on SPY confirms the gap: Black-Scholes delta remains competitive on real crash windows when training is simulation-based.

Phase III: Data Aligned Recovery. Backtest training closes much of the 2008 in-distribution gap, showing that distribution shift, not network capacity, dominates the first real-market failure. Expanded-corpus training falsifies a simple regime-coverage story; COVID remains hard because daily rebalancing with a lagging vega proxy is the wrong tool for gap risk.

Phase IV: Instrument and Architecture Fix. VIX futures provide forward-looking vega with leverage-corrected training. Multi-seed analysis reveals a domain of validity: robust calm-market wins, unstable crisis losses. The regime-gated architecture is the first design that closes the COVID gap without sacrificing the 2024 out-of-training win, with seed standard deviation reduced by an order of magnitude.

Scientific stance. The final claim is intentionally narrower than the initial thesis. The framework is demonstrably strong in simulation and calm real markets, and its crisis failure modes are now diagnosed rather than concealed. This transparency is a feature of the report; it defines precisely where the method is deployable and what remains open for future investigation.

26 Replication Guidelines

This section maps each major result block to its entry-point script, expected outputs, and reproduction order. Full command-level detail is also provided in `REPLICATION.md` in the GitHub repository.

26.1 Environment Setup

1. Clone `Deep-Hedging-Cvar-Heston` and create a Python 3.10+ virtual environment.
2. Install dependencies: `pip install -r requirements.txt`.
3. Run all commands from the repository root so that `results/` paths resolve correctly.
4. Optional GPU: CUDA accelerates training; CPU is sufficient for several evaluation-only scripts.



26.2 Experiment to Script Map

Table 16: Replication map: report sections to runnable scripts.

Report section	Entry script	Key outputs
Statistical robustness	<code>main.bootstrap.py</code>	<code>bootstrap_ci_plot.png</code>
Jump diffusion	<code>main.bates.py</code>	<code>bates_vs_heston.png</code>
SPX calibration	<code>main.calibrated.py</code> , <code>main.spx_calibration.py</code>	<code>spx_calib_comparison.png</code>
Heston benchmark	<code>main.heston_benchmark.py</code>	<code>heston_benchmark.png</code>
OTM / portfolio / learning	<code>main.otm_sweep.py</code> , <code>main.portfolio_hedge.py</code> , <code>main.learning_curve.py</code>	<code>matching_results/*.png</code>
Risk comparison	<code>main.risk_comparison.py</code>	<code>risk_comparison_cvar.png</code>
Misspecification	<code>main.misspecification.py</code>	<code>misspecification_cvar.png</code>
Domain randomisation	<code>main.robust_training.py</code>	<code>robust_training_cvar.png</code>
Greeks decomposition	<code>main.greeks_decomposition.py</code>	<code>greeks_decomposition_waterfall.png</code>
Multi-asset spread	<code>main.multi_asset.py</code>	<code>results/multi_asset/*.png</code>
Ablations	<code>main.ablation_mean_cvar.py</code> , <code>main.ablation_design.py</code>	<code>results/ablation/*.png</code>
Historical stress	<code>main.historical_stress.py</code>	<code>results/historical_stress_cvar.png</code>
Sim-to-real	<code>main.sim_to_real.py</code>	metrics table in stdout
Backtest training	<code>main.backtest_training.py</code>	<code>results/backtest_training/*.png</code>
Expanded corpus	<code>main.expanded_corpus.py</code>	<code>results/expanded_corpus_cvar.png</code>
VIX futures	<code>main.vix_futures.py</code> , <code>main.vix_futures_v2.py</code>	<code>results/vix_futures_v2_cvar.png</code>
Multi-seed verification	<code>main.vix_multiseed.py</code>	<code>results/vix_multiseed_delta.png</code>
Regime-gated architecture	<code>main.vix_gated_multiseed.py</code>	gated metrics in stdout
Gate threshold validation	<code>main.gate_threshold_validation.py</code>	<code>results/report/gate_threshold_sweep.png</code>

26.3 Recommended Replication Order

1. **Baseline pipeline:** `python main.py` then `python main.varswap_cvar.py`.
2. **Core simulated evidence:** bootstrap, Bates, benchmark, misspecification, robust training, Greeks decomposition.
3. **Extensions:** multi-asset, ablations, path-dependent payoffs.
4. **Real-market arc:** historical stress → sim-to-real → backtest training → expanded corpus → VIX futures v2 → multi-seed → gated multi-seed.
5. **Figure regeneration:** `python make_report_plots.py` for consolidated report figures where applicable.

26.4 Data Dependencies

- Simulated experiments require no external data beyond random seeds.
- SPX calibration uses option quotes via the calibration module (`calibration/calibrate_spx.py`).
- Real-market scripts download SPY and VIX via `yfinance` on first run and cache windows under `data/`.
- Reproducing exact floating-point results may require matching PyTorch/CUDA versions; seed-level qualitative conclusions are stable.



26.5 Minimum Replication Checklist

To verify the three central thesis claims without rerunning the full suite:

1. Claim (i): `main_heston_benchmark.py`: stock+VS beats analytical ceiling.
2. Claim (ii): `main_portfolio_hedge.py` and `main_path_dependent.py`: extensions remain positive.
3. Claim (iii): `main_misspecification.py` then `main_robust_training.py`: failure and partial cure.
4. Real-market capstone: `main_vix_gated_multiseed.py` with frozen gate $\theta = 3.0$.

27 Conclusion and Future Work

27.1 Conclusion

This report consolidates a year-long research programme on risk-sensitive deep hedging with variance instruments, conducted under the Dual Degree Project framework at IIT Madras. The central simulated result stands robustly: a CVaR-trained policy augmented with a variance swap instrument can outperform both stock-only deep hedging and the instantaneous analytical Heston delta-vega ceiling at realistic transaction costs. This finding holds across Bates jump-diffusion paths, SPX-calibrated parameter sets, path-dependent payoffs, option portfolios, and multi-asset spread options, the most generality achieved by any single study combining stochastic volatility, CVaR objectives, and variance instruments in the published deep hedging literature.

The real-market programme is an equally important contribution. Simulation-only training does not dominate Black-Scholes delta on crisis windows, an honest negative result documented fully and without equivocation. Backtest-based training resolves much of the in-distribution gap; expanded-corpus training falsifies a naive regime-coverage explanation; VIX futures with leverage correction and multi-seed verification characterise a clean calm-market domain of validity. The regime-gated architecture represents the first design in this project to simultaneously preserve the calm-market outperformance and close the COVID crisis gap with low cross-seed variance, satisfying both pre-registered decision rules.

The final scientific position is precise and deliberately bounded: the framework constitutes a strong and well-understood hedging learner under controlled stochastic-volatility environments, and a promising, rigorously diagnosed approach on calm real-market regimes. Crisis robustness requires explicit regime conditioning, not more simulation alone. The boundaries of validity are not weaknesses to be concealed; they are results that define a principled and deployable research contribution.

27.2 Future Work

1. **Neural SDE simulator.** Replace hand-specified Heston sampling with learned generative dynamics (`main_nsde_train.py`, `main_nsde_hedge.py`).
2. **Principled gate selection.** Hold out a validation period before final OOT evaluation (`main_gate_threshold_v`).
3. **Tradable microstructure.** Incorporate realistic VIX futures costs and liquidity constraints (`main_vix_with_costs.py`).



4. **Sequence policies.** LSTM-gated architectures for path dependence in vega control (`main_lstm_gate.py`).
5. **Portfolio-scale deployment.** Extend from single-option hedging to book-level CVaR with cross-asset netting.

References

- [1] Buehler, H., Gonon, L., Teichmann, J., & Wood, B. (2019). Deep hedging. *Quantitative Finance*, 19(8), 1271–1291.
- [2] Horvath, B., Muguruza, A., & Tomas, M. (2021). Deep learning volatility. *Quantitative Finance*, 21(1), 11–27.
- [3] Carbonneau, A., & Godin, F. (2021). Equal risk pricing of derivatives with deep hedging. *Quantitative Finance*, 21(4), 593–608.
- [4] Imaki, S., Jimba, M., & Nakajima, H. (2023). No-transaction band network. *Journal of Financial Data Science*, 5(2), 84–103.
- [5] Marzban, S., Delage, E., & Li, J. Y.-M. (2022). Deep reinforcement learning for equal risk pricing and hedging. *Quantitative Finance*, 22(9), 1641–1664.
- [6] Bates, D. S. (1996). Jumps and stochastic volatility. *Review of Financial Studies*, 9(1), 69–107.
- [7] Heston, S. L. (1993). A closed-form solution for options with stochastic volatility. *Review of Financial Studies*, 6(2), 327–343.
- [8] Gatheral, J., & Jacquier, A. (2011). Convergence of Heston to SVI. *Quantitative Finance*, 11(8), 1129–1132.
- [9] Rockafellar, R. T., & Uryasev, S. (2000). Optimization of conditional value-at-risk. *Journal of Risk*, 2, 21–42.
- [10] Margrabe, W. (1978). The value of an option to exchange one asset for another. *Journal of Finance*, 33(1), 177–186.
- [11] Pinto, L., Davidson, J., Sukthankar, R., & Gupta, A. (2017). Robust adversarial reinforcement learning. *ICML*.
- [12] Politis, D. N., & Romano, J. P. (1994). The stationary bootstrap. *Journal of the American Statistical Association*, 89(428), 1303–1313.
- [13] Lugosi, G., & Mendelson, S. (2019). Mean estimation and regression under heavy-tailed distributions. *Foundations of Computational Mathematics*, 19(5), 1145–1190.
- [14] Kloeden, P. E., & Platen, E. (1992). *Numerical Solution of Stochastic Differential Equations*. Springer.

A Hyperparameter Reference

Table 17 lists all fixed Heston simulation parameters used in the base experiments. Table 18 lists the shared neural network and training hyperparameters.



Table 17: Heston model parameters used in simulation (base configuration).

Parameter	Symbol	Value
Mean reversion speed	κ	3.0
Long-run variance	θ	0.04
Vol-of-vol	σ_v	0.3
Correlation	ρ	-0.7
Initial variance	v_0	0.04
Initial spot	S_0	100.0
CVaR level	α	0.05
Risk-free rate	r	0.0
Drift	μ	0.0

Table 18: Neural network and training hyperparameters (shared across all experiments).

Hyperparameter	Value
Architecture	3-layer MLP, hidden size 64, ReLU
Input features	$(S_t, v_t, t/T, \text{moneyness})$
Output	Action logits (1 or 2 instruments)
Final activation	$\text{Tanh} \times \text{max-delta clip (1.0)}$
Optimiser	Adam ($\beta_1 = 0.9, \beta_2 = 0.999$)
Learning rate	10^{-3}
LR schedule	CosineAnnealing, $T_{\max} = 200$
Batch size	256 paths
Epochs	200
CVaR estimator	Soft-plus approximation, $\epsilon = 10^{-3}$
Domain rand. support	$\kappa \in [1, 6], \theta \in [0.01, 0.09], \sigma_v \in [0.1, 0.6], \rho \in [-0.9, -0.3]$
Seeds	5 (results averaged, std reported)
Path steps	$T = 100$
Paths per epoch	$N = 10\,000$
Transaction cost	$\lambda = 0.001$ (0.1%)

Regime Gate Parameters

The sigmoid gate used in the Regime Gated experiment (Section 23) was calibrated to the following fixed values and *not* re-optimised during neural network training:



Table 19: Fixed parameters of the VIX-based regime gate.

Parameter	Symbol	Value
Gate centre (VIX ratio)	θ_{gate}	3.0
Gate width	w	0.3
VS weight scale	—	0.3
VIX normalisation baseline	V_0	V_{open} (session start)

These parameters were set prior to any experiment run and were not tuned on test data, in line with the pre-registration requirement.